

```

In [1]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import load_digits
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import classification_report, accuracy_score

# Load inbuilt Digits dataset
digits = load_digits()
X = digits.data # Each image is 8x8 pixels, flattened to 64 features
y = digits.target

# Split into training, validation, and test sets
X_train, X_temp, y_train, y_temp = train_test_split(X, y, test_size=0.3, random_state=42)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42)

# Standardize features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_val = scaler.transform(X_val)
X_test = scaler.transform(X_test)

# Train MLP model
mlp = MLPClassifier(hidden_layer_sizes=(128,), activation='relu', solver='adam', max_iter=50, random_state=42)
mlp.fit(X_train, y_train)

# Evaluate model
val_preds = mlp.predict(X_val)
test_preds = mlp.predict(X_test)

print("Validation Accuracy:", accuracy_score(y_val, val_preds))
print("Test Accuracy:", accuracy_score(y_test, test_preds))
print("\nClassification Report:\n", classification_report(y_test, test_preds))

# Plot training loss curve
plt.plot(mlp.loss_curve_)
plt.title("Training Loss Curve")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.grid(True)
plt.show()

# Display sample predictions
def display_predictions(images, labels, model):
    plt.figure(figsize=(10, 4))
    for i in range(5):
        plt.subplot(2, 5, i + 1)
        plt.imshow(images[i].reshape(8, 8), cmap='gray')
        plt.title(f"Label: {labels[i]}")
        plt.axis('off')

        plt.subplot(2, 5, i + 6)
        probs = model.predict_proba(images[i].reshape(1, -1))[0]
        plt.bar(range(10), probs)
        plt.title(f"Pred: {np.argmax(probs)}")
        plt.xticks(range(10))
    plt.tight_layout()
    plt.show()

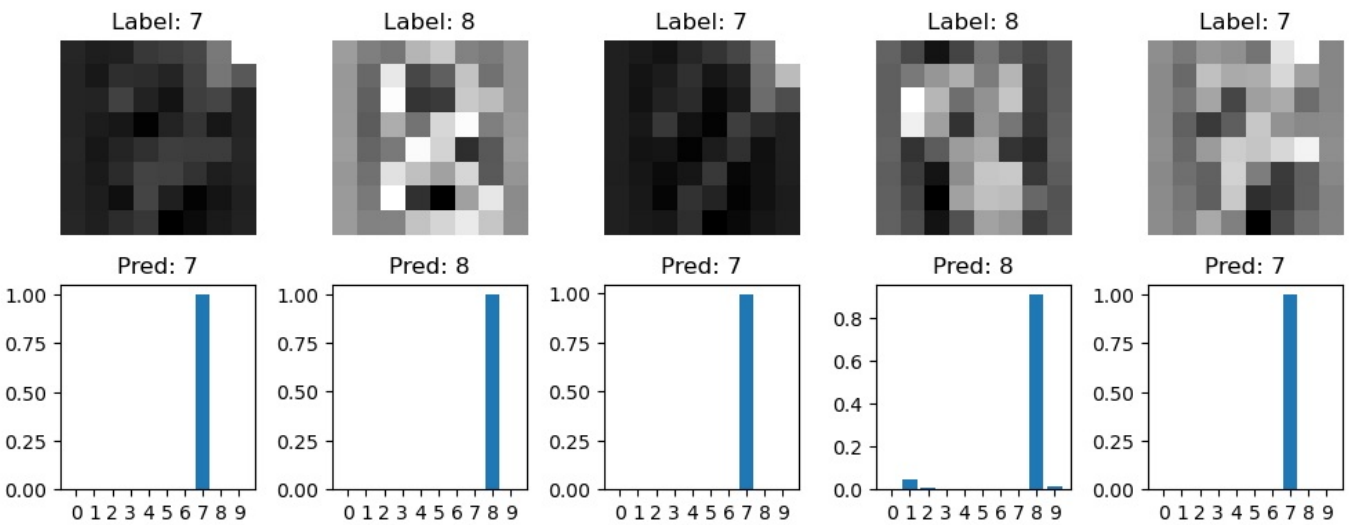

# Show predictions on random test samples
indices = np.random.choice(len(X_test), 5, replace=False)
display_predictions(X_test[indices], y_test[indices], mlp)

```

Validation Accuracy: 0.9740740740741
Test Accuracy: 0.9740740740741

Classification Report:				
	precision	recall	f1-score	support
0	1.00	0.97	0.98	31
1	1.00	0.96	0.98	24
2	0.93	1.00	0.96	27
3	0.97	0.97	0.97	33
4	0.94	1.00	0.97	29
5	0.96	0.93	0.94	27
6	1.00	0.96	0.98	23
7	1.00	1.00	1.00	27
8	1.00	1.00	1.00	25
9	0.96	0.96	0.96	24
accuracy			0.97	270
macro avg	0.98	0.97	0.97	270
weighted avg	0.97	0.97	0.97	270

C:\Users\User\anaconda3\Lib\site-packages\sklearn\normal_multilayer_perceptron.py:691: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (50) reached and the optimization hasn't converged yet.
warnings.warn()



In []: